

# I/O System Calls

## Chapter 3

modified from slides by Dr. B. Boufama and Dr. Quazi Rahman

## Unix I/O System Calls

- most Unix I/O can be performed using **system calls**
  - open, close
  - read, write
  - lseek
- in contrast to **standard I/O functions**, they are **unbuffered**

## Comparison

I/O system calls	standard I/O library
unbuffered I/O	buffering
not part of ISO C	specified in C standard

## File Descriptors

- who use it
  - kernel refers to any open file by a *file descriptor*
- what is it
  - non-negative integer **0..OPEN\_MAX**
- how do you get it
  - typically returned by system call `open()` to open/create these
    - file
    - pipe
    - network communication channel

## File Descriptors

<unistd.h>

	file descriptors	value
standard input	STDIN_FILENO	0
standard output	STDOUT_FILENO	1
standard error	STDERR_FILENO	2

1/2/19

COMP-2560 System Programming

University of Windsor

5

## System Call: open()

```
#include <fcntl.h>
```

```
int open(const char *fName, int oflag, .../* mode_t mode */);
```

returns: file descriptor if ok, -1 on error

- *fName*
  - name of file to open
- *oflag*
  - list of values separated by bitwise 'OR' (|)
  - determines how the file is to be opened
  - its value defined in <fcntl.h> (file control)

1/2/19

COMP-2560 System Programming

University of Windsor

6

## System Call: open()

- value for *oflag*
  - must have one of these three

oflag value	meaning
O_RDONLY	read only
O_WRONLY	write only
O_RDWR	read and write

## System Call: open()

- value for *oflag*
  - these are optional

oflag	meaning...
O_APPEND	open for writing at end of file
O_CREAT	<ul style="list-style-type: none"> <li>• if file does not exist, create it</li> <li>• if this option is used, also need third parameter (<i>mode_t mode</i>) for access permission bits</li> </ul>
O_EXCL	<ul style="list-style-type: none"> <li>• generate error if O_CREAT is specified and file already exists</li> <li>• ensure that caller must create the file</li> </ul>
O_TRUNC	if file exists and successfully opened for either write-only or read-only, truncate its length to zero

## System Call: open()

- value for *mode* defined in <fcntl.h>

	mode mask	bit mask
access permission for owner: R, W, X	S_IRUSR	0400
	S_IWUSR	0200
	S_IXUSR	0100
access permission for group: R, W, X	S_IRGRP	0040
	S_IWGRP	0020
	S_IXGRP	0010
access permission for others: R, W, X	S_IROTH	0004
	S_IWOTH	0002
	S_IXOTH	0001

## System Call: creat()

```
#include <fcntl.h>
int creat(const char *fName, mode_t mode);
```

returns: file descriptor if ok, -1 o.w.

equivalent to

```
open(fName, O_WRONLY | O_CREAT | O_TRUNC, mode)
```

## System Call: close()

```
#include <unistd.h>
int close(int fd);
```

returns: 0 if ok, -1 on error

- what happens with this operator
  - file descriptor is set free
- when to return -1 ?
  - e.g. *fd* was already closed
- what happens when process terminates
  - all open files closed by kernel

## System Call: lseek()

```
#include <unistd.h>
off_t lseek(int fd, off_t offset, int whence);
```

returns: new file offset if ok, -1 on error

- what does it do
  - set *file pointer* associated with file *fd*

## System Call: lseek()

### *offset*

- a long integer
- can be a negative number
- meaning: number of bytes
- initialized to 0 when file opened, unless O\_APPEND is specified

## System Call: lseek()

### *whence*

<i>whence</i> is SEEK_SET	set file offset to <i>offset</i> bytes from beginning
<i>whence</i> is SEEK_CUR	set file offset to current value plus <i>offset</i>
<i>whence</i> is SEEK_END	set file offset to size of the file plus <i>offset</i>

## System Call: lseek()

how to use lseek:

- use *lseek* to determine the current offset

off\_t currpos

currpos = lseek(fd, 0, SEEK\_CUR)

## System Call: lseek()

how to use lseek:

- use *lseek* to determine if a file is capable of seeking
  - *lseek* returns -1 if *fd* refers to pipe, socket, etc. (cannot seek)



## System Call: lseek()

```
#include <unistd.h>
```

```
int main(void) {
```

```
    if (lseek(STDIN_FILENO, 0, SEEK_CUR)  
    == -1)
```

```
        printf("cannot seek\n");
```

```
    else
```

```
        printf("seek ok\n");
```

```
    exit(0);
```

```
}
```

pipe or socket

```
> a.out < testfile  
seek ok  
> cat testfile | a.out  
cannot seek
```

1/2/19

COMP-2560 System Programming

University of Windsor

17

## System Call: lseek()

*offset*

– what happens if offset is greater than file's current size ?

- file gets extended
- any bytes not written are read back as 0

Example...  
pay attention to file  
size and file content

1/2/19

COMP-2560 System Programming

University of Windsor

18

```

#include <unistd.h>
#include <fcntl.h>
#include <stdlib.h>

int main() {

    char buf1[] = "abcdefghij";
    char buf2[] = "ABCDEFGHIJ";
    int fd;

    if (( fd = creat("newfile", 0533 )) < 0) {
        printf("creat error\n");
        exit(1);
    }
    if (( write(fd, buf1, 10) ) != 10) {
        printf("write error\n");
        exit(1);
    }
    if (( lseek(fd, 15, SEEK_SET) ) == -1) {
        printf("write error\n");
        exit(1);
    }
    if (( write(fd, buf2, 10) ) != 10) {
        printf("write error\n");
        exit(1);
    }
    close(fd);
}

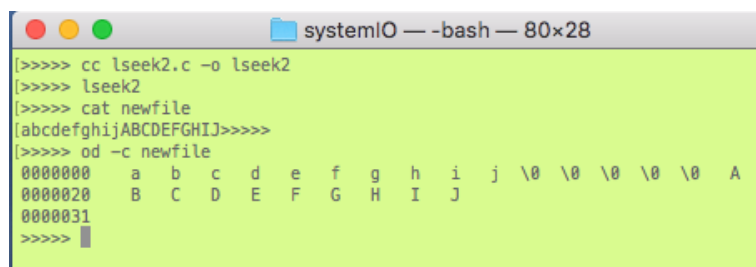
```

1/2/19

COMP-2560 System Programming

University of Windsor

19



```

systemIO — -bash — 80x28
[>>>> cc lseek2.c -o lseek2
[>>>> lseek2
[>>>> cat newfile
[abcdefghijABCDEFGHIJ]>>>>
[>>>> od -c newfile
0000000  a  b  c  d  e  f  g  h  i  j  \0 \0 \0 \0 \0  A
0000020  B  C  D  E  F  G  H  I  J
0000031
>>>> █

```

1/2/19

COMP-2560 System Programming

University of Windsor

20

## System Call: read()

```
#include <unistd.h>
ssize_t read(int fd, void *buf, size_t nbytes);
```

returns: number of bytes read,  
0 if end of file  
-1 on error

- read up to *nbytes*
- *ssize\_t* usually defined as *signed size\_t*
- in what situations: *return value = 0* ?  
(EOF has already been reached)

## System Call: read()

```
#include <unistd.h>
ssize_t read(int fd, void *buf, size_t nbytes);
```

returns: number of bytes read,  
0 if end of file  
-1 on error

- in what situations: *return value < nbytes* ?
  - number of bytes read before EOF is less than *nbytes*
  - reading from keyboard – up to a line
  - reading from a network

```

Terminal — ssh — 80x29
xjchen@charlie:~/60256/demo$ ls
cprogram.c  myscript  text
xjchen@charlie:~/60256/demo$ cat cprogram.c
#include <stdio.h>
#include <unistd.h>
#include <fcntl.h>
#include <stdlib.h>

int main() {

    int buf = 'a';
    read(STDIN_FILENO, &buf, 1);
    printf("%d\n", buf);
    read(STDIN_FILENO, &buf, 1);
    printf("%d\n", buf);
}
xjchen@charlie:~/60256/demo$ cc cprogram.c
xjchen@charlie:~/60256/demo$ a.out
x
120
10
xjchen@charlie:~/60256/demo$

```

user typed: "x" followed by "return"

1/2/19

COMP-2560 System Programming

University of Windsor

23

## System Call: write()

```
#include <unistd.h>
```

```
ssize_t write(int fd, const void *buf, size_t nbytes);
```

returns: nbytes if ok, -1 o.w.

1/2/19

COMP-2560 System Programming

University of Windsor

24

## Example: reversing a file

- what does it do
  - last byte becomes the first byte
  - second last byte becomes the second
  - ...
- command
  - > reverse fileIn fileOut
- two solutions
  - use lseek on fileIn (read bytes from end)
  - use lseek on fileOut (write bytes from end)

1/2/19

COMP-2560 System Programming

University of Windsor

25

solution 1: use lseek() on fileIn

```
#include <unistd.h>
#include <fcntl.h>
int main(int argc, char *argv[]) {
    int fd1, fd2;
    char buffer;
    long int i = 0, fileSize = 0;

    fd1 = open(argv[1], O_RDONLY);
    fd2 = open(argv[2], O_CREAT|O_WRONLY|O_TRUNC, 0755);

    while (read(fd1, &buffer, 1) > 0)
        fileSize++;

    while (++i <= fileSize) {
        lseek(fd1, -i, SEEK_END);
        read(fd1, &buffer, 1);
        write(fd2, &buffer, 1);
    }
    close(fd1);
    close(fd2);
}
```

1/2/19

COMP-2560 System Programming

University of Windsor

26

```
//solution 2: start writing at the end

int main(int argc, char *argv[]) {

    int fd1, fd2;
    char buffer;
    long int i = 0, fileSize = 0;

    fd1 = open(argv[1], O_RDONLY);
    fd2 = open(argv[2], O_CREAT|O_WRONLY|O_TRUNC, 0755);

    while(read(fd1, &buffer, 1)>0)
        fileSize++;

    lseek(fd2, fileSize-1, SEEK_SET);           // extend
    lseek(fd1, 0, SEEK_SET);

    while(++i <= fileSize) {
        read(fd1, &buffer, 1);
        lseek(fd2, -i, SEEK_END);
        write(fd2, &buffer, 1);
    }
    close(fd1);
    close(fd2);
}
```