

Process Control

Chapter 8

modified from slides by Dr. B. Boufama and Dr. Quazi Rahman

Unix process

create new
process: fork

terminate
process: exit

waiting for a child
process: wait()/waitpid()

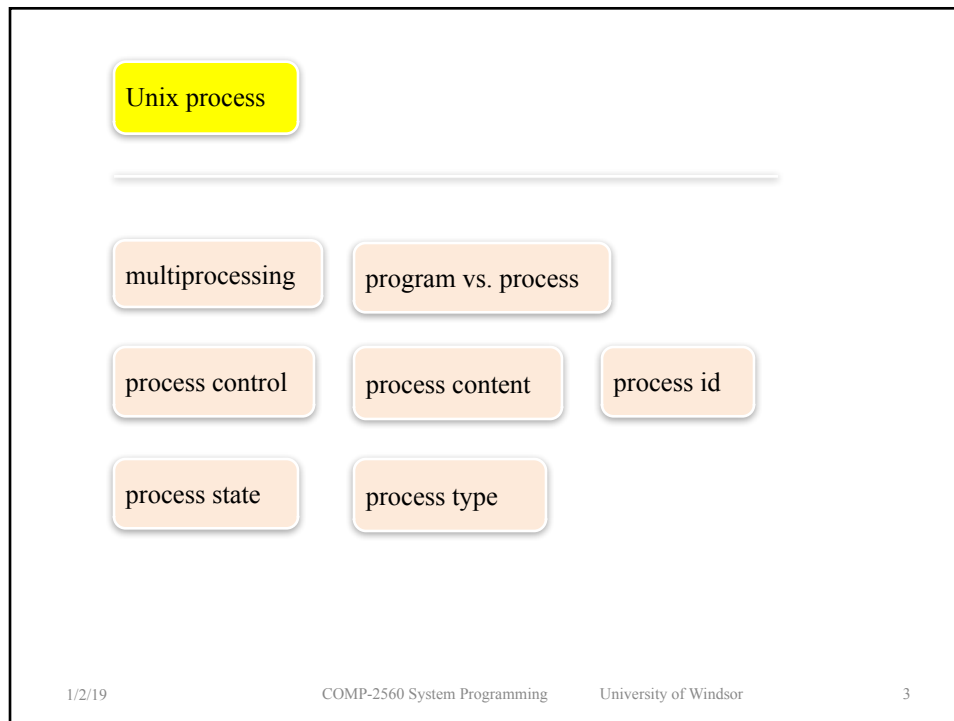
orphan and
zombie process

differentiating
process: exec()

changing
directory: chdir()

system scheduling
priority

process group



Multiprocess

- multiprocessing
 - simultaneously executing programs on same machine
 - illusion: executions done in parallel
- multiprogramming
 - CPU switches among programs
 - illusion: all programs continuously executing

Program vs. Process

- program
 - (an executable) file (residing on a disk)
- process
 - executing (running) program
 - opened in the working memory (RAM)
 - usually with a limited life-time
 - also called *task*
- a running program -> process

1/2/19

COMP-2560 System Programming

University of Windsor

5

Process Control

- what is included in process control
 - creation of new processes
 - program execution
 - process termination
- what is included in a process
 - process ID
 - user ID, group ID

1/2/19

COMP-2560 System Programming

University of Windsor

6

What Does a Unix Process Contain ?

- a unique **process ID**
- **user ID** of the owner
- **code segment**
- data segment (variables)
- **stack segment**
- **an environment**

- nonnegative
- assigned by OS
- used to identify a process

instructions that are being executed

a form of memory where it is possible to push and pop instructions

e.g.
registers' contents
tables of open files

Process ID

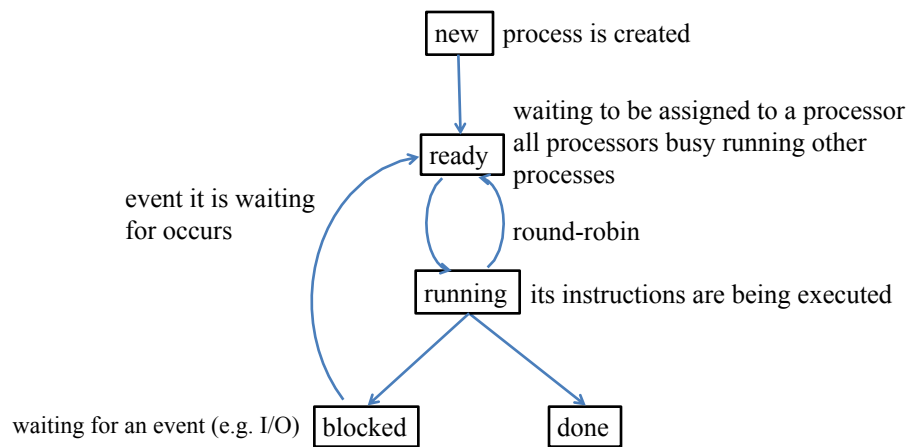
- system calls **getpid()** and **getppid()**

```
#include <stdio.h>

int main(void) {
    printf("Hello, my PID is %d\n", getpid());
    printf("Hello, my PPID is %d\n", getppid());
    exit(0);
}
```

```
> a.out
Hello, my PID is 11723
Hello, my PPID is 5598
```

Process State



1/2/19

COMP-2560 System Programming

University of Windsor

9

Process Types

process types	examples
system process	scheduler process
user process	init process

1/2/19

COMP-2560 System Programming

University of Windsor

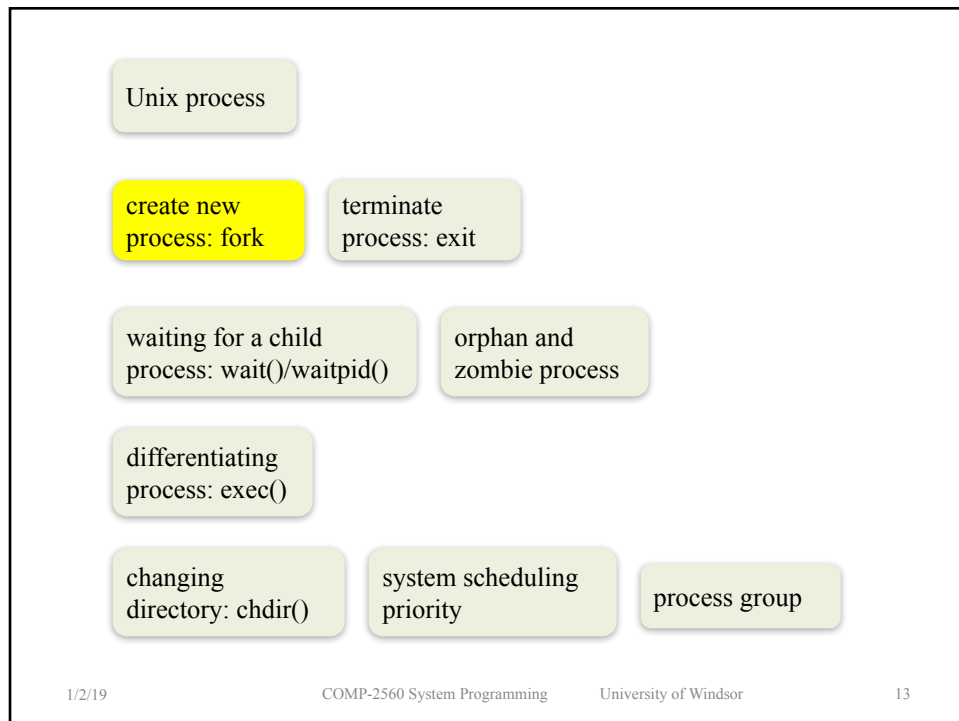
10

System Process vs. User Process

- scheduler process (PID=0)
 - no program on disk corresponds to this process
 - part of kernel

System Process vs. User Process

- init process (PID=1)
 - invoked by kernel at the end of bootstrap procedure
 - normal user process, not a system process
 - program file is on disk (e.g. /etc/init)
 - read system-dependent initialization files
 - never dies: continues running until system shut down



Create New Process

- how to create a new process
 - duplicate + replace
(either in two steps or in one step)
 - distinguish two processes: child vs. parent
 - *init* is ancestor of all subsequent processes

Create New Process

- system calls
 - `fork()`: duplicate caller process
 - `exec()`: replace the caller process by a new one
 - *spawn*: a single operation for *fork* followed by *exec*

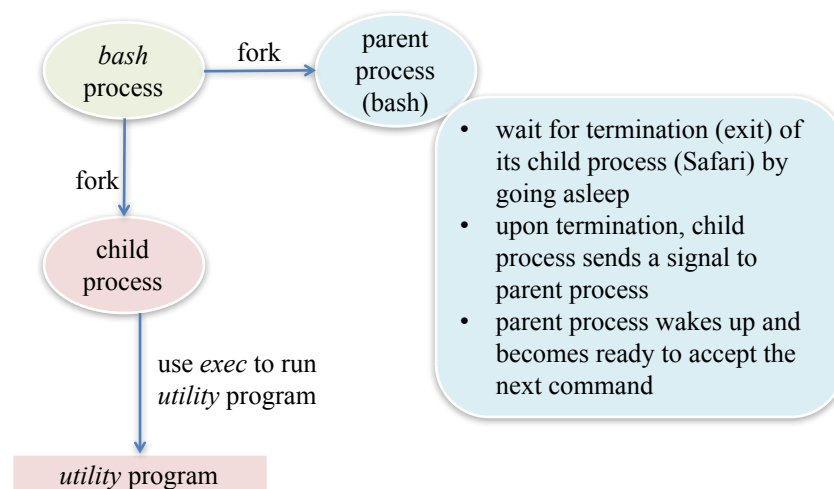
1/2/19

COMP-2560 System Programming

University of Windsor

15

Example 2: Running Safari from Shell



1/2/19

COMP-2560 System Programming

University of Windsor

16

System Call fork()

- synopsis

`pid_t fork(void)`

- what does it do?
 - try to duplicate caller process
- what does it return?
 - if successful ?
 - if not successful ?

1/2/19

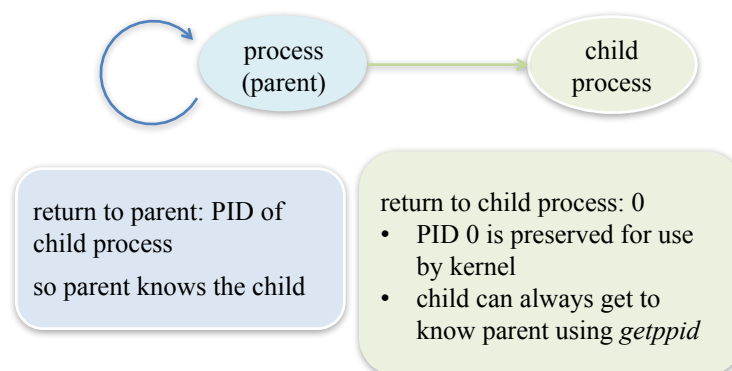
COMP-2560 System Programming

University of Windsor

17

System Call fork()

if successful



1/2/19


COMP-2560 System Programming

University of Windsor

18

System Call fork()

if not successful



two main reasons fork() fails

- too many processes already in system
- total number of processes for this real user ID exceeds system's limit

CHILD_MAX specifies maximum number of simultaneous processes per real user ID

System Call fork()

- fork() is a special (strange) system call
 - called once by one process
 - return twice, to two different processes
- a child process contains
 - its own PID
 - its parent process ID
 - its own copy of the parent's data segment and file descriptors

System Call fork()

- both parent and child resume execution
 - who starts execution first ?

who is executing?

```
#include <unistd.h>
int main(int argc, char *argv[]){
    int npid;
    printf("Initially, PID = %d\n", getpid());
    npid = fork();
    if(npid == -1) {
        perror("impossible to fork");
        exit(1);
    }
    printf("my npid =%d, my PID =%d\n", npid, getpid())
    exit(0);
}
```

```
#include <unistd.h>
```

fork() – data segment

```
int glob = 100;
```

```
int main(){
```

```
    int pid,
```

```
    int var = 88;
```

```
    printf("pid=%d, glob=%d, var=%d\n", getpid(), glob, var);
```

```
    pid = fork();
```

```
    if ( pid < 0 )    exit(1);
```

```
    if (pid == 0){
```

```
        glob++;
```

```
        var++;
```

```
    }
```

```
    else
```

```
        sleep(2);
```

```
    printf("pid=%d, glob=%d, var=%d\n", getpid(), glob, var);
```

```
    exit(0);
```

```
}
```

1/2/19

COMP-2560 System Programming

University of Windsor

23

Example: fork() – data segment

```
>>>>> a.out
```

```
pid = 430, glob = 101, var = 89
```

```
pid = 429, glob = 100, var = 88
```

1/2/19

COMP-2560 System Programming

University of Windsor

24

fork(): file descriptor

- child has parent's file descriptors
- parent and child share file offset
- if parent and child both write to same descriptor
 - output will be intermixed
- see example...

1/2/19

COMP-2560 System Programming

University of Windsor

25

```
#include <fcntl.h>
#include <unistd.h>
int main(){
    int pid, fd, i;
    char c;
    if ( (fd = open("test", O_RDWR | O_CREAT, 0700)) == -1 ) {
        perror("failed to open test"); exit(0);
    }
    if( (pid = fork()) < 0 ) {
        perror(" failed to fork"); exit(1);
    }
    if(pid == 0)
        for( i = 65; i < 85; i++) {
            c = i;
            write(fd, &c, 1); // child print ABCDEFGHIJKLMNOPQRST
        }
    else {
        sleep(1);
        for( i = 0 ; i < 20; i++) {
            c = 58;
            write(fd, &c, 1);      /*character : = 58 */
        }
    }
    return 0;
}
```

ABCDEFGHIJKLMNOPQRSTUVWXYZ:.....

26

exercise: what is output?

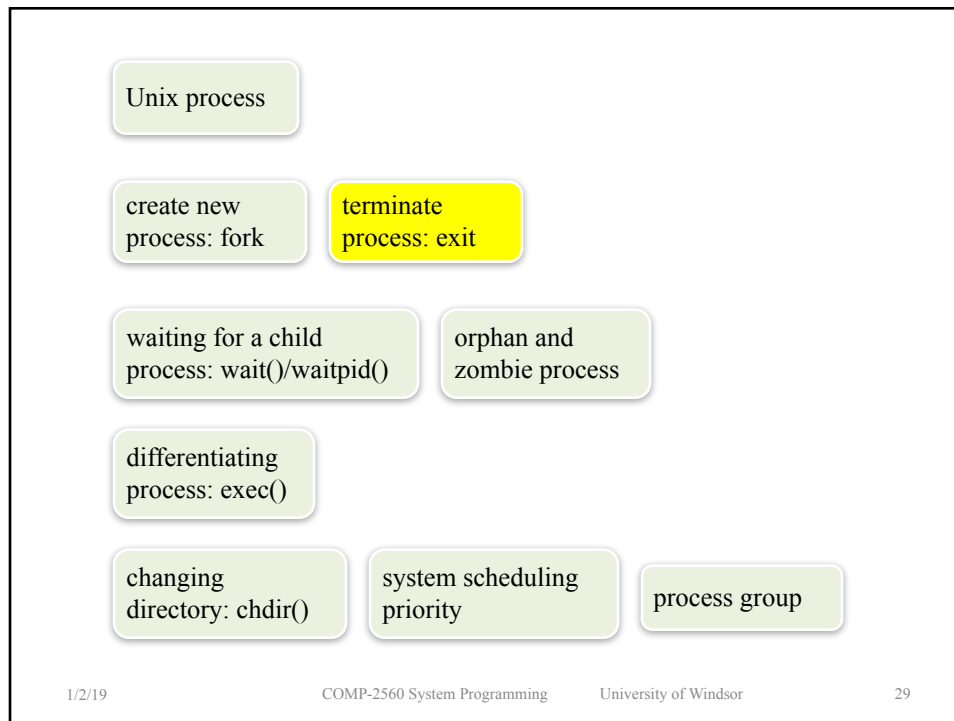
```
int main(){
    fork();
    fork();
    fork();
    printf("done, my pid is %d\n", getpid());
}
```

```
done, my pid is 15958
done, my pid is 15959
done, my pid is 15962
done, my pid is 15963
done, my pid is 15965
done, my pid is 15961
done, my pid is 15964
done, my pid is 15960
```

exercise: what is output?

```
int main(){
    int i;
    printf("Before fork, my pid is %d\n", getpid());
    for (i=0; i<3; i++){
        if ( fork()== 0 )
            printf("Hi, I am child. My pid is %d\n", getpid());
    }
}
```

```
Before fork, my pid is 3163
Hi, I am child. My pid is 3164
Hi, I am child. My pid is 3165
Hi, I am child. My pid is 3166
Hi, I am child. My pid is 3169
Hi, I am child. My pid is 3168
Hi, I am child. My pid is 3167
Hi, I am child. My pid is 3170
```



Process Termination: `exit()`

- synopsis


```
void exit(int status);
```
- terminate a process and never return
- what does it do
 - close all file descriptors
 - flush all output streams and close all open streams
 - free memory used by its code, data, stack
 - send a SIGCHLD signal to its parent and wait for parent to accept its returned code

`exit()`: never returns

```
int main() {
    int newpid;
    printf("before: my pid is %d\n", getpid());
    if ((newpid = fork()) == -1 )
        perror("fork");
    else if (newpid == 0) {
        printf("I am the child %d now sleeping...\n", getpid());
        sleep(1);
        exit(47);
        printf("I am gone");
    }
    else {
        printf("I am the parent %d\n", getpid());
        sleep(3);
        printf("My child %d must be gone by now. I am leaving...\n", newpid);
        exit(1);
        printf("I am gone too\n");
    }
}
```

31

`exit()` - discussions

- what if parent terminates before child?
 - *init* process becomes parent

when a process terminates, kernel goes through all active process to check

- change parent process ID of surviving processes to 1

exit() - discussions

- problem when child terminates before parent
 - parent may want to check termination status of a child
 - termination status lost when child disappears
 - kernel keeps minimal info about terminating child for its parent
 - process ID
 - termination status
 - etc.

1/2/19

COMP-2560 System Programming

University of Windsor

33

Unix process

create new
process: fork

terminate
process: exit

waiting for a child
process: wait()/waitpid()

orphan and
zombie process

differentiating
process: exec()

changing
directory: chdir()

system scheduling
priority

process group

1/2/19

COMP-2560 System Programming

University of Windsor

34

wait()

```
#include <sys/wait.h>
pid_t wait(int *statloc);
        return: process ID if ok, 0 or -1 on error
```

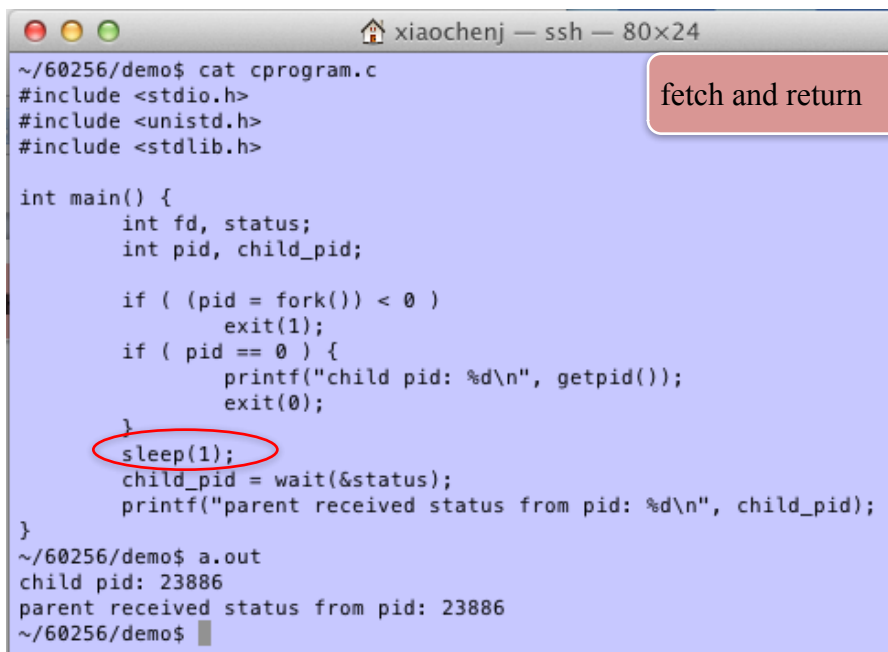
- what does *wait* do
 - if all children are running, get blocked
 - if no child process, return with error
 - if a child terminated and waiting for its termination status to be fetched, return with status

1/2/19

COMP-2560 System Programming

University of Windsor

35



```
~/60256/demo$ cat cprogram.c
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>

int main() {
    int fd, status;
    int pid, child_pid;

    if ( (pid = fork()) < 0 )
        exit(1);
    if ( pid == 0 ) {
        printf("child pid: %d\n", getpid());
        exit(0);
    }
    sleep(1);
    child_pid = wait(&status);
    printf("parent received status from pid: %d\n", child_pid);
}

~/60256/demo$ a.out
child pid: 23886
parent received status from pid: 23886
~/60256/demo$
```

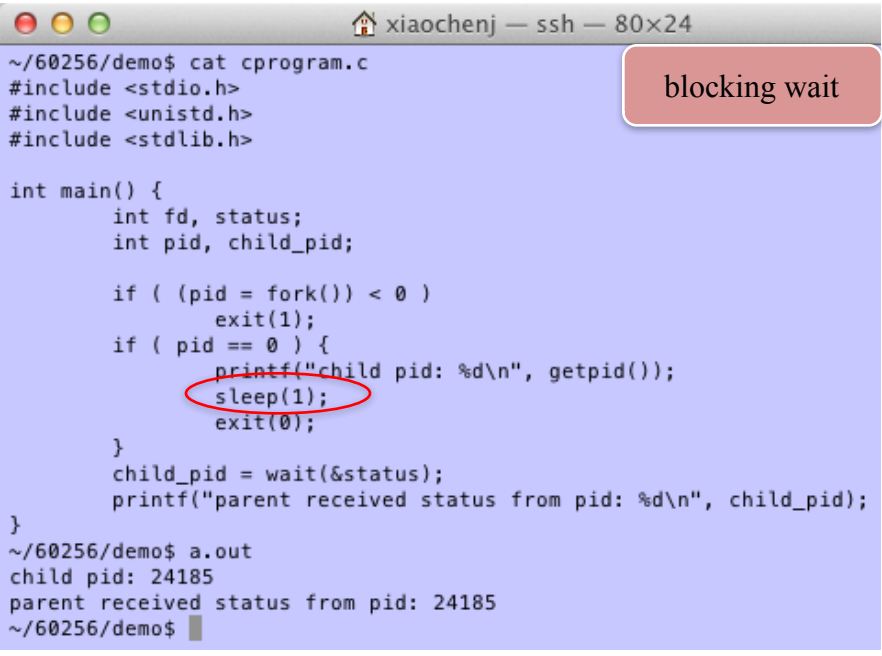
fetch and return

1/2/19

COMP-2560 System Programming

University of Windsor

36



```

~/60256/demo$ cat cprogram.c
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>

int main() {
    int fd, status;
    int pid, child_pid;

    if ( (pid = fork()) < 0 )
        exit(1);
    if ( pid == 0 ) {
        printf("child pid: %d\n", getpid());
        sleep(1);
        exit(0);
    }
    child_pid = wait(&status);
    printf("parent received status from pid: %d\n", child_pid);
}

~/60256/demo$ a.out
child pid: 24185
parent received status from pid: 24185
~/60256/demo$

```

blocking wait

1/2/19 COMP-2560 System Programming University of Windsor 37

wait()

```
#include <sys/wait.h>
pid_t wait(int *statloc);
```

return: process ID if ok, 0 or -1 on error

- with return PID value of *wait*, we can tell which child terminated

calling *wait* when there are more than one child process

```

#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main() {
    int status;
    pid_t pid;

    pid = fork();
    if ( pid == -1 ) perror("fork");
    if ( pid == 0 ) {
        printf("first child: %d\n", getpid());
        sleep(50);
        exit(0);
    }
    pid = fork();
    if ( pid == -1 ) perror("fork");
    if ( pid == 0 ) {
        printf("second child: %d\n", getpid());
        exit(0);
    }
    printf("return from wait: %d\n", wait(&status));
}

```

```

first child: 490
second child: 491
return from wait: 491

```

1/2/19

COMP-2560 System Programming


University of Windsor

39

wait()

- *statloc* – a pointer to an integer

- the leftmost byte contains the status returned by child
- it is a value 0-255 (passed as an argument to exit)
- represent normal termination of child



if rightmost byte is zero

1/2/19

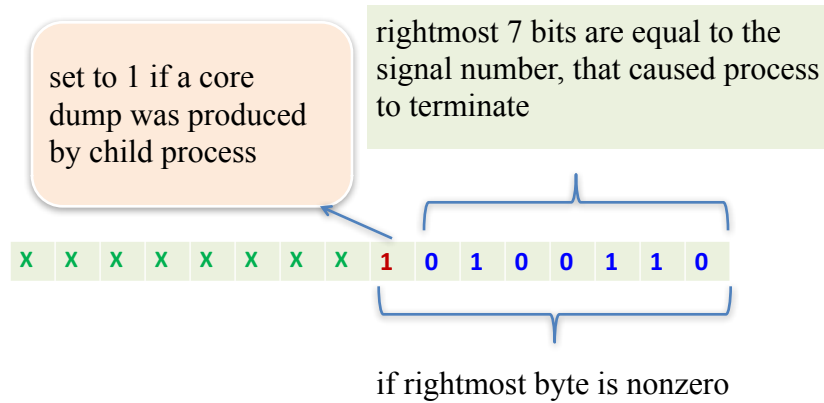
COMP-2560 System Programming

University of Windsor

40

wait()

- *statloc* – a pointer to an integer



1/2/19

COMP-2560 System Programming

University of Windsor

41

wait(): core dump

- core dump
 - refers to a file (named core)
 - consists of recorded state of working memory at a specific time
- usually recorded when program terminated abnormally (crashed)
- often used to assist in diagnosing and debugging errors

1/2/19

COMP-2560 System Programming

University of Windsor

42

Example

- let child process `exit with 47`
- the parent will have status value: 2f00 (12032)

abort()

- synopsis

`void abort(void)`

- declared in `<stdlib.h>`
- causes abnormal process termination to occur
- it sends signal SIGABRT (6) to the parent process
- causes a core dump

Example

- above example again
- let child process call `abort()`
- the parent will have status value: `0086` (SIGABRT 6, with core dump)
- to enable core dump on your machine, you can set

> `ulimit -c unlimited`

1/2/19

COMP-2560 System Programming

University of Windsor

45

```

>>>> more abort86.c
#include <stdio.h>
#include <fcntl.h>
#include <stdlib.h>
#include <unistd.h>

int main() {
    int fd, status;
    pid_t pid = fork();
    if ( pid < 0 )
        exit(1);
    if ( pid == 0 )
        abort();
    wait(&status);

    printf("status: %x\n", status);
    printf("status: %d\n", status);

    fd = open("datafile", O_CREAT | O_TRUNC | O_WRONLY, 0700);
    write(fd, &status, 2);
    close(fd);
}

>>>> cc abort86.c ; a.out
status: 86
status: 134
>>>> xxd datafile
00000000: 8600                                     ..
>>>>

```

1/2/19

COMP-2560 System Programming

University of Windsor

46

Example

- now, before the child makes exit call, let us **terminate the child process**
 > kill -15 child-id
- the parent will have status value: **000f**
 (SIGTERM 15)

child process: sleep and get external signal to terminate

```

>>>> more kill15.c
#include <stdio.h>
#include <fcntl.h>
#include <stdlib.h>
#include <unistd.h>

int main() {
    int fd, status;
    pid_t pid = fork();
    if ( pid < 0 )
        exit(1);
    if ( pid == 0 ) {
        printf("child process id: %d\n", getpid());
        sleep(50);
        exit(47);
    }
    printf("parent process id: %d\n", getpid());
    wait(&status);

    printf("status: %x\n", status);
    printf("status: %d\n", status);

    fd = open("datafile", O_CREAT | O_TRUNC | O_WRONLY, 0700);
    write(fd, &status, 2);
    close(fd);
}

>>>>

```


Retrieving Status Information

- include `<sys/wait.h>`

WIFEXITED(status) (if exited)	true for normal child termination
WEXITSTATUS(status) (exit status)	returns exit status as an integer (0-255) used only when WIFEXITED(status) is true
WIFSIGNALED(status) (if signaled)	true for abnormal child termination
WTERMSIG(status)	returns signal number that caused abnormal child death used only when WIFSIGNALED(status) is true
WCOREDUMP(status)	true if a core file was generated

1/2/19

COMP-2560 System Programming

University of Windsor

49

previous example again, with WIFSIGNALED

```
int main() {
    int status;
    pid_t pid;
    pid = fork();
    if ( pid == -1 )
        perror("fork");
    if ( pid == 0 ) {
        printf("child process id: %d\n", getpid());
        sleep(50);
        exit(47);
    }
    else {
        printf("parent process id: %d\n", getpid());
        wait(&status);
        if ( WIFSIGNALED(status) )
            printf("signal number: %d\n", WTERMSIG(status));
    }
}
```

waitpid()

```
pid_t waitpid(pid_t pid, int *status, int options)
```

- wait for a *specific* child process
- return *error* if
 - specified process does not exist
 - specified process group does not exist
 - specified process is not a child of calling process

waitpid()

- argument *pid*

passing argument pid	meaning
pid == -1	wait for any child (equivalent to <i>wait</i>)
pid > 0	wait for child with <i>pid</i>
pid == 0	wait for any child whose group ID equals that of calling process
pid < -1	wait for any child whose group ID is absolute value of <i>pid</i>

`wait(&status)` is equivalent to `waitpid(-1, &status, 0)`

waitpid()

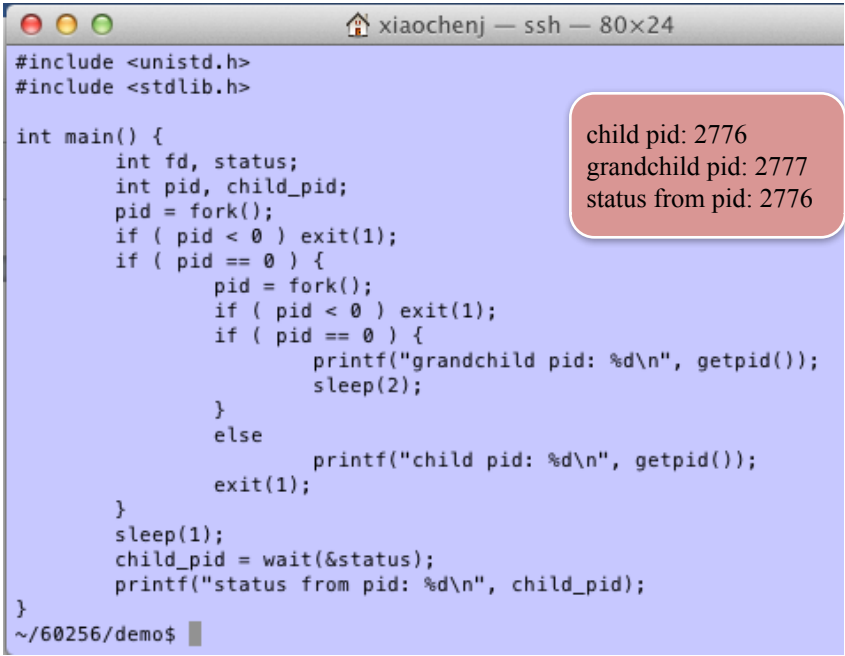
```
pid_t waitpid(pid_t pid, int *status, int options)
```

- argument *option*
 - can be 0
 - can be e.g. **WNOHANG**
 - specify nonblocking: the waitpid() will not block if a child specified by *pid* is not immediately available
 - in this case, the return value is 0

non-blocking
version of wait()

waitpid()

- note:
 - waitpid() waits for a child process,
grandchildren not counted*
- first example: child exits first, grandchild second
- second example: grandchild exits first, child second



```

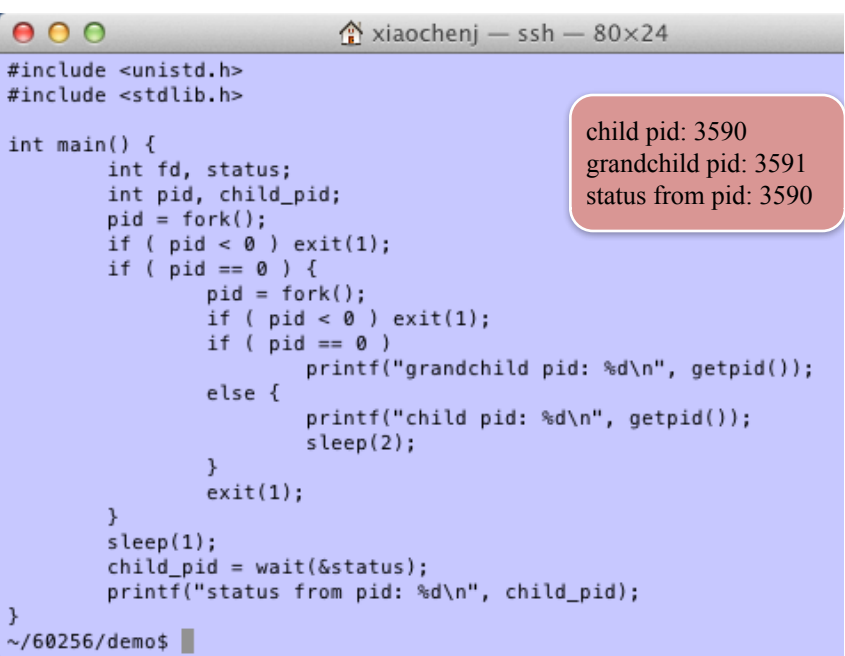
#include <unistd.h>
#include <stdlib.h>

int main() {
    int fd, status;
    int pid, child_pid;
    pid = fork();
    if ( pid < 0 ) exit(1);
    if ( pid == 0 ) {
        pid = fork();
        if ( pid < 0 ) exit(1);
        if ( pid == 0 ) {
            printf("grandchild pid: %d\n", getpid());
            sleep(2);
        }
        else
            printf("child pid: %d\n", getpid());
        exit(1);
    }
    sleep(1);
    child_pid = wait(&status);
    printf("status from pid: %d\n", child_pid);
}
~/60256/demo$

```

child pid: 2776
grandchild pid: 2777
status from pid: 2776

1/2/19 COMP-2560 System Programming University of Windsor 55



```

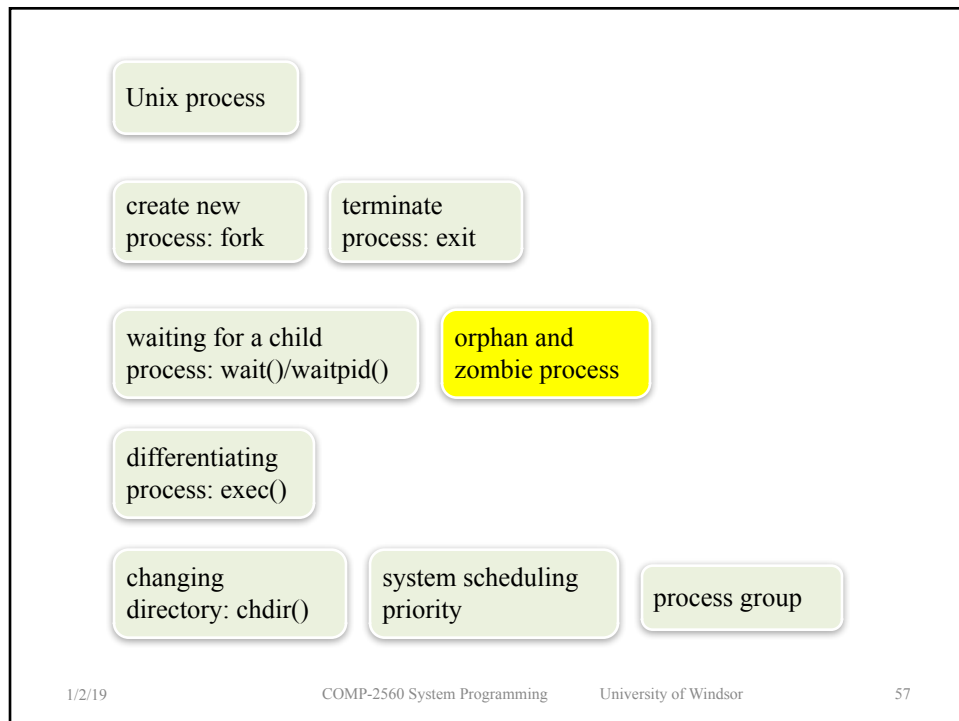
#include <unistd.h>
#include <stdlib.h>

int main() {
    int fd, status;
    int pid, child_pid;
    pid = fork();
    if ( pid < 0 ) exit(1);
    if ( pid == 0 ) {
        pid = fork();
        if ( pid < 0 ) exit(1);
        if ( pid == 0 ) {
            printf("grandchild pid: %d\n", getpid());
        }
        else {
            printf("child pid: %d\n", getpid());
            sleep(2);
        }
        exit(1);
    }
    sleep(1);
    child_pid = wait(&status);
    printf("status from pid: %d\n", child_pid);
}
~/60256/demo$

```

child pid: 3590
grandchild pid: 3591
status from pid: 3590

1/2/19 COMP-2560 System Programming University of Windsor 56



Orphan and Zombie

- zombie process
 - a terminated process
 - its parent has not yet waited for it
- what if a child of *init* terminates?
 - it does not become zombie
 - *init* always calls one of *wait* functions to fetch its termination status

Orphan and Zombie

- terminated process does not leave system before its parent accepts its return

- | | |
|--|--|
| <ul style="list-style-type: none">• parent exits<ul style="list-style-type: none">– e.g. parent killed prematurely• child alive<ul style="list-style-type: none">– become <i>orphan</i> | <ul style="list-style-type: none">• parent alive but no call to <code>wait()</code>• child terminated<ul style="list-style-type: none">– become <i>zombie</i> |
|--|--|

Orphan

- orphan processes are systematically adopted by *init*
 - kernel changes PPID of orphan to 1

becoming child of *init* process

```

Terminal — ssh — 80x24
~/60256/demo$ cat cprogram.c
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>

int main() {
    pid_t pid;
    if (( pid = fork() ) < 0 ) {
        perror("fork");
        exit(1);
    }
    if ( pid > 0 ) {
        sleep(1);
        exit(0);
    }
    printf("parent pid is: %d\n", getppid());
    sleep(2);
    printf("parent pid is: %d\n", getppid());
}
~/60256/demo$ a.out
parent pid is: 15178
~/60256/demo$ parent pid is: 1
~/60256/demo$

```

1/2/19

COMP-2560 System Programming

University of Windsor

61

Zombie process

- zombies
 - compared to normal processes, they lose their resources e.g.
 - data
 - code
 - stack
 - however, they remain in system's process table waiting for acceptance of their return (system's process table has a fixed size)

1/2/19

COMP-2560 System Programming

University of Windsor

62

Example: making a zombie process

```

Terminal — ssh — 80x24
~/60256/demo$ cat cprogram.c
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>

int main() {
    int pid;
    if ( (pid = fork()) < 0 )
        exit(1);
    if ( pid > 0 ) {
        printf("parent pid: %d\n", getpid());
        while (1)
            sleep(5);
    }
    printf("child pid: %d\n", getpid());
}
~/60256/demo$

```

1/2/19

COMP-2560 System Programming

University of Windsor

63

command to check the running processes

> `ps -u your_user_id`

```

xiaochenj — ssh —
~/60256/demo$ a.out
parent pid: 4961
child pid: 4962
^Z
[1]+  Stopped                  a.out
~/60256/demo$ ps -u xjchen
  PID TTY          TIME CMD
 4271 ?            00:00:00 sshd
 4272 pts/19      00:00:00 bash
 4961 pts/19      00:00:00 a.out
 4962 pts/19      00:00:00 a.out <defunct>
 5001 pts/19      00:00:00 ps
23158 ?            00:00:00 sshd
23159 pts/9       00:00:00 bash
~/60256/demo$

```

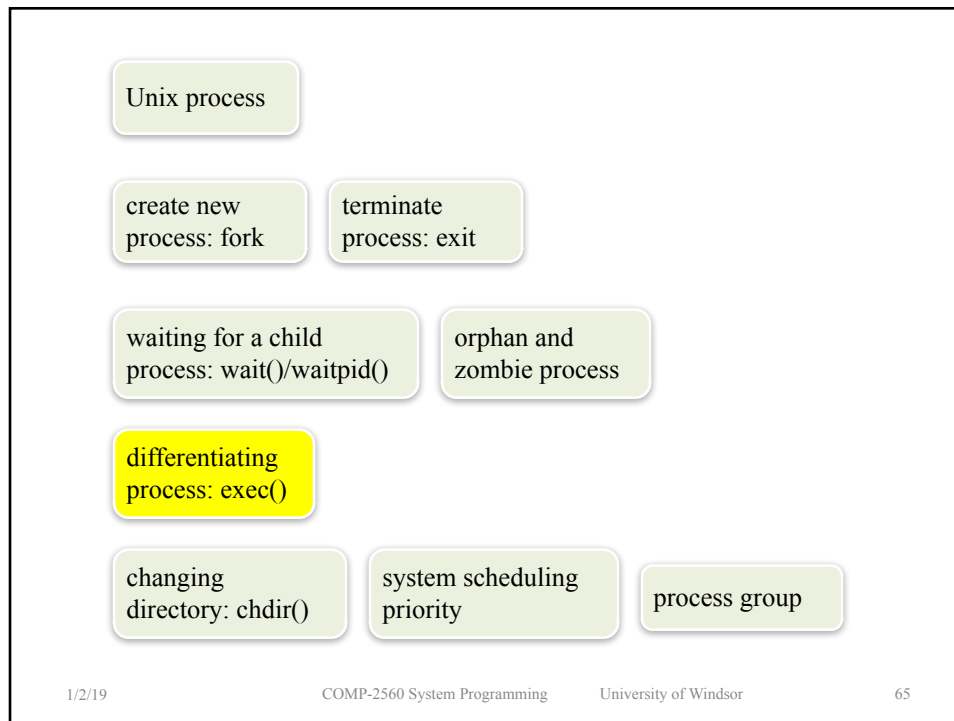
to resume (after Control-Z),
you can use *bg* or *fg*

1/2/19

COMP-2560 System Programming

University of Windsor

64



Differentiating Process: exec()

```
# include <unistd.h>
int execl(const char *pathname, const char *arg0, ..., (char *)0);
int execv(const char *pathname, char *const argv[]);
int execlp(const char *pathname, const char *arg0, ..., (char *)0);
int execvp(const char *pathname, char *const argv[]);

return: -1 on error, no return on success
```

- exec() family of system calls
- replace current process
 - code
 - data
 - stack

Differentiating Process: exec()

```
# include <unistd.h>
int execl(const char *pathname, const char *arg0, ..., (char *)0);
int execv(const char *pathname, char *const argv[]);
int execlp(const char *pathname, const char *arg0, ..., (char *)0);
int execvp(const char *pathname, char *const argv[]);

return: -1 on error, no return on success
```

- new program starts executing *main* function
- process ID not changed
 - no new process created
- a successful call to exec() never returns (why?)

1/2/19

COMP-2560 System Programming

University of Windsor

67

Differentiating Process: exec()

	first parameter	other parameters
execl	complete pathname of the executable	<ul style="list-style-type: none"> • arg0 must be program name • the list of arguments must be terminated by a NULL pointer
execlp	use \$PATH to find program	
execv	complete pathname of the executable	<ul style="list-style-type: none"> • arg0 must be program name • the array of pointers must be terminated by a NULL pointer
execvp	use \$PATH to find program	

1/2/19

COMP-2560 System Programming

University of Windsor

68

```

int main(int argc, char* argv[]) {
    int pid;
    printf("Before fork: process id %d\n", getpid());
    if ((pid = fork()) == 0) {
        printf("I am the child %d\n", getpid());
        sleep(5);
        printf("Listing content of current directory...\n");
        execl("/bin/ls", "ls", "-l", (char *)0);
    }
    else {
        printf("I am the parent %d\n", getpid());
        int status, term_pid = wait(&status);
        printf("Child %d listed the content of current directory\n", term_pid);
        exit(1);
    }
}

```

1/2/19

COMP-2560 System Programming

University of Windsor

69

```

int main(int argc, char* argv[]) {
    int pid;
    printf("Before fork: process id %d\n", getpid());
    if ((pid = fork()) == 0) {
        printf("I am the child %d\n", getpid());
        sleep(5);
        printf("Listing content of current directory...\n");
        execlp("ls", "ls", "-l", 0);
    }
    else{
        printf("I am the parent %d\n", getpid());
        int status;
        int term_pid = wait(&status);
        printf("Child %d has listed the content of current directory\n", term_pid);
        exit(1);
    }
}

```

1/2/19

COMP-2560 System Programming

University of Windsor

70

```

int main(int argc, char* argv[]) {
    int pid;
    printf("Before: process id %d\n", getpid());
    if ((pid = fork())==0) {
        printf("I am the child %d\n", getpid());
        sleep(5);
        printf("Listing content of current directory...\n");
        char* arg_list[3] = {"ls", "-l", (char *)0};
        execv("/bin/ls", arg_list);
    }
    else {
        printf("I am the parent %d\n", getpid());
        int status, term_pid = wait(&status);
        printf("Child %d has listed the content of current directory\n", term_pid);
        exit(1);
    }
}

```

or
char* arg_list[3];
arg_list[0] = "ls";
arg_list[1] = "-l";
arg_list[2] = 0;

1/2/19

COMP-2560 System Programming

University of Windsor

71

```

int main(int argc, char* argv[]) {
    int pid;
    printf("Before: process id %d\n", getpid());
    if ((pid = fork())==0) {
        printf("I am the child %d\n", getpid());
        sleep(5);
        printf("Listing content of current directory...\n");
        char* arg_list[3] = {"ls", "-l", (char *)0};
        execvp("ls", arg_list);
    }
    else {
        printf("I am the parent %d\n", getpid());
        int status, term_pid = wait(&status);
        printf("Child %d has listed the content of current directory\n", term_pid);
        exit(1);
    }
}

```

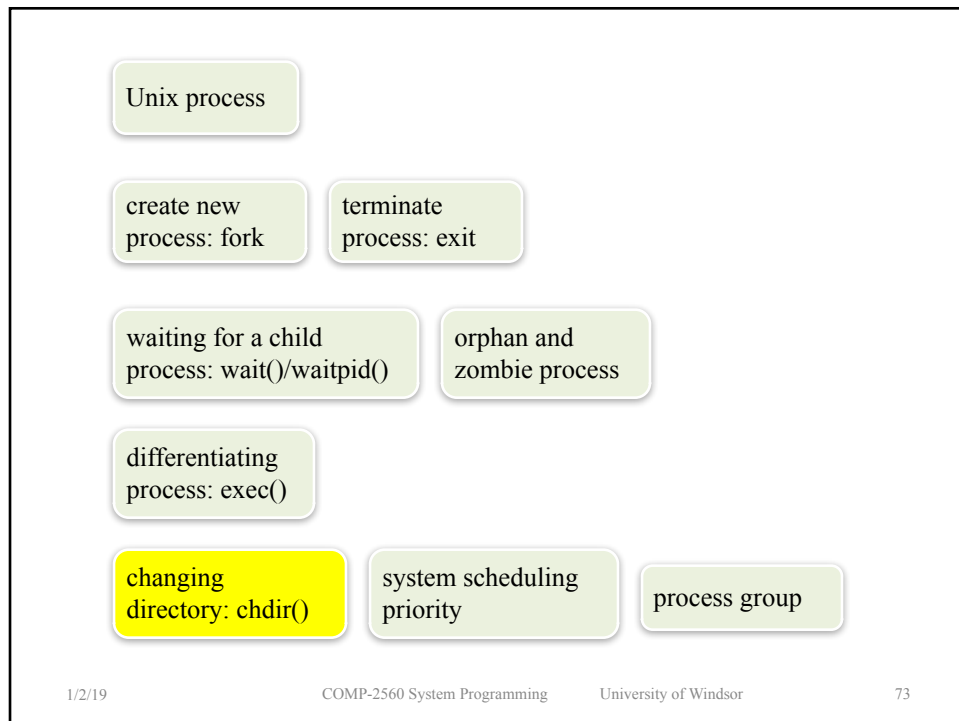
or
char* arg_list[3];
arg_list[0] = "ls";
arg_list[1] = "-l";
arg_list[2] = 0;

1/2/19

COMP-2560 System Programming

University of Windsor

72



Inheriting and Changing Directory

- child process inherits current working directory of parent
- child process can change working directory using *chdir()*

Inheriting and Changing Directory

```
int chdir(const char *pathname);
```

- return 0 if successful
- return -1 if fails
 - specified path name does not exist
 - the process does not have execute permission from the directory

1/2/19

COMP-2560 System Programming

University of Windsor

75

Unix process

create new
process: fork

terminate
process: exit

waiting for a child
process: wait()/waitpid()

orphan and
zombie process

differentiating
process: exec()

changing
directory: chdir()

system scheduling
priority

process group

1/2/19

COMP-2560 System Programming

University of Windsor

76

System Scheduling Priority

- each process has a system scheduling priority
 - each process runs at a default system priority: 0
 - child priority inherited from its parent
- priority value range
 - -20 ~ 19
 - range differs from one Unix platform to another
 - negative values restricted to super-user

1/2/19

COMP-2560 System Programming

University of Windsor

77

System Scheduling Priority

- priority values affect amount of CPU time allocated to the process
 - smaller the value, faster the process
- changing scheduling priority
 - nice()
 - setpriority()

1/2/19

COMP-2560 System Programming

University of Windsor

78

getpriority() and setpriority()

<sys/resource.h>

```
int getpriority(int which, id_t who)
```

returns priority value on success; -1 on failure

- obtains current scheduling priority of a process, process group, or user
- *which*
 - identifies whether it is a process (PRIO_PROCESS), process group, etc.
- *who*
 - process id, group id etc
 - if *who* is 0, the calling process (or group etc.) is considered

1/2/19

COMP-2560 System Programming

University of Windsor

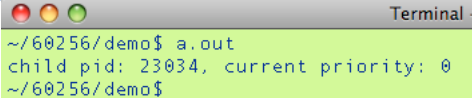
79

Example: getpriority()

```
#include <stdio.h>
#include <unistd.h>
#include <sys/wait.h>
#include <stdlib.h>

int main() {
    int pid, status;
    if ( (pid = fork()) < 0 )
        exit(1);
    if ( pid == 0 )
        printf("child pid: %d, current priority: %d\n",
               getpid(), getpriority(PRIO_PROCESS, getpid()));

    else
        wait(&status);
}
```



```
Terminal
~/60256/demo$ a.out
child pid: 23034, current priority: 0
~/60256/demo$
```

1/2/19

COMP-2560 System Programming

University of Windsor

80

System Scheduling Priority

```
int nice(int delta)
```

- adds delta to current value
- only super-user processes can have a negative value
- returns new priority value if successful; -1 o.w.

1/2/19

COMP-2560 System Programming

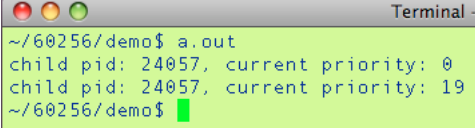
University of Windsor

81

example: nice()

```
#include <stdio.h>
#include <unistd.h>
#include <sys/wait.h>
#include <stdlib.h>

int main() {
    int pid, status;
    if ( (pid = fork()) < 0 )
        exit(1);
    if ( pid == 0 ) {
        printf("child pid: %d, current priority: %d\n",
               getpid(), getpriority(PRIO_PROCESS, getpid()));
        nice(19);
        printf("child pid: %d, current priority: %d\n",
               getpid(), getpriority(PRIO_PROCESS, getpid()));
    }
    else
        wait(&status);
}
```



```
Terminal
~/60256/demo$ a.out
child pid: 24057, current priority: 0
child pid: 24057, current priority: 19
~/60256/demo$
```

1/2/19

COMP-2560 System Programming

University of Windsor

82

getpriority() and setpriority()

<sys/resource.h>

int setpriority(int **which**, id_t **who**, int **priority**)

returns 0 on success; -1 on failure

1/2/19

COMP-2560 System Programming

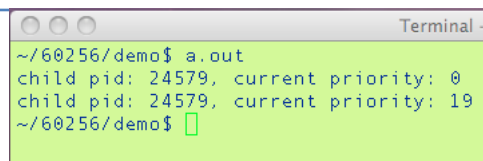
University of Windsor

83

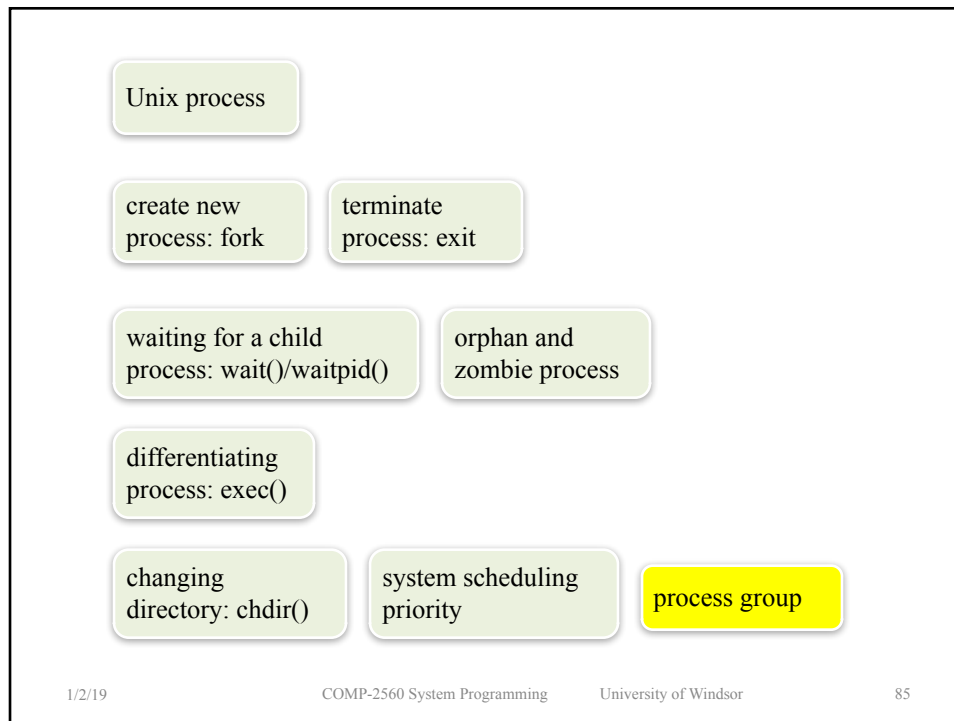
example: parent sets priority of the child

```
#include <stdio.h>
#include <unistd.h>
#include <sys/wait.h>
#include <stdlib.h>

int main() {
    int pid, status;
    if ( (pid = fork()) < 0 )
        exit(1);
    if ( pid == 0 ) {
        printf("child pid: %d, current priority: %d\n",
               getpid(), getpriority(PRIO_PROCESS, getpid()));
        sleep(5);
        printf("child pid: %d, current priority: %d\n",
               getpid(), getpriority(PRIO_PROCESS, getpid()));
    }
    else {
        sleep(1);
        setpriority(PRIO_PROCESS, pid, 19);
        wait(&status);
    }
}
```



```
Terminal —
~/60256/demo$ a.out
child pid: 24579, current priority: 0
child pid: 24579, current priority: 19
~/60256/demo$
```



Process Groups

- every process is a member of a process group
 - a child inherits process group from parent
 - when calling exec(), process group remains the same
- one of group members is the group leader
 - each group member has the group leader's process ID and its process-group-ID

Process Groups

- a process may change its process group
 - to another group
 - create its own group
 - being leader and sole member
- kernel provides a system call to send a signal to each member of a designated process group
 - can be used to terminate the entire group

1/2/19

COMP-2560 System Programming

University of Windsor

87

setpgid() and getpgid()

```
<sys/types.h>
<unistd.h>
int setpgid(pid_t pid, pid_t pgid)
```

- set process group of process with *pid* to *pgid*
- returns 0 if successful; -1 o.w.
- if *pgid* == *pid*, the process becomes process group leader
- if *pid* == 0, process ID of the calling process is used
- if *pgid* == 0, the process ID *pid* is used
 - process specified by *pid* becomes a process group leader

1/2/19

COMP-2560 System Programming

University of Windsor

88

setpgid() and getpgid()

```
<sys/types.h>
<unistd.h>

pid_t getpgid(pid_t pid)
```

- return process group id of the process with *pid*
- if *pid* == 0, the calling process group ID is returned

Example: parent id and group id

```
#include <unistd.h>
#include <sys/wait.h>

int main(int argc, char *argv[]){

    printf("Parent: PID = %d, PPID = %d, PGID = %d\n",
           getpid(), getppid(), getpgid(getpid()));

    if ( fork() == 0 ) {
        printf("Child: PID = %d, PPID = %d, PGID = %d\n",
               getpid(), getppid(), getpgid(getpid()));
        exit(1);
    }
    sleep(5);
}
```

Output:

```
Parent: PID = 20814, PPID = 20381, PGID = 20814
Child: PID = 20815, PPID = 20814, PGID = 20814
```

example: set group id

```

#include <unistd.h>
#include <sys/wait.h>

int main(int argc, char *argv[]){

    printf("Parent: PID = %d, PPID = %d, PGID = %d\n",
           getpid(), getppid(), getpgid(getpid()));

    if ( fork() == 0 ){
        printf("Child: PID = %d, PPID = %d, PGID = %d\n",
               getpid(), getppid(), getpgid(getpid()));

        setpgid( 0, 0 );           //or setpgid(getpid(),0);

        printf("Child after setpgid: PID = %d, PPID = %d, PGID = %d\n",
               getpid(), getppid(), getpgid(getpid()));
    }
    sleep(5);
}

```

Output:

```

Parent: PID = 22295, PPID = 20381, PGID = 22295
Child: PID = 22296, PPID = 22295, PGID = 22295
Child after setpgid: PID = 22296, PPID = 22295, PGID = 22296

```